


Forward Rendering

Drawing things one at a time

Render Pipeline So Far...

- Everything we've discussed so far is about how the GPU works when processing data
- How the hardware works is one thing, how you make use of it is a completely different story
- The way in which your application renders its information is the most important factor when writing graphics-based applications

001011011010001010102



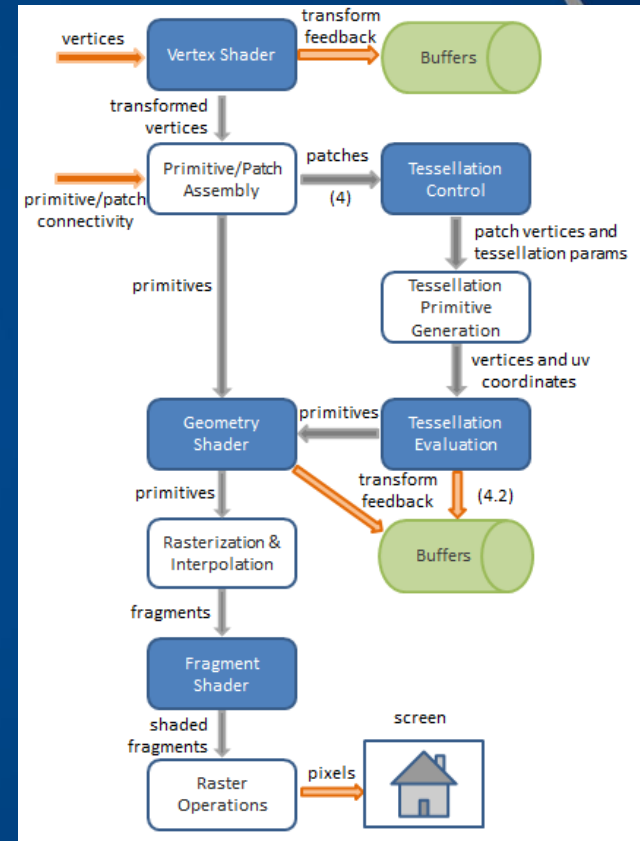
workplace • Online • Campus | cit.edu.au

my of
five
inment
cit.edu.au




Render Steps

- So recapping the steps in rendering:
 - Start a draw call for a collection of primitives
 - Vertex Shader processes every vertex in the mesh independently from each other
 - The Tessellators processes all patches if it needs
 - The Geometry Shader then processes every primitive independently from each other
 - The Rasteriser then determines if a pixel should be rendered
 - The Fragment Shader then processes all pixels independently from each other
 - The Raster Operations perform any additional processing on the pixels
 - Repeat all steps for the next mesh being rendered...
- This process is referred to as Forward Rendering
 - In a future session we will discuss its opposite; Deferred Rendering




Ordering – The Painter's Algorithm

- One big problem exists in rendering
 - Visibility
- If I render a mesh, then render another mesh over it, which one should be visible?
 - One solution to this problem is called The Painter's Algorithm
- The Painter's Algorithm dictates that a scene should be drawn back to front, so that no distant object overlaps a closer object




Painter's Algorithm Flaws

- Although being a great way to solve the issue, the algorithm still suffers from a flaw
 - How would it handle the following geometric arrangement?
- Also it means that thousands of pixels are being rendered multiple times in a single frame, even though only a single pixel colour will be visible in the end
 - In the previous example the pixels the trees are rendered into were already rendered as grass, which was a waste



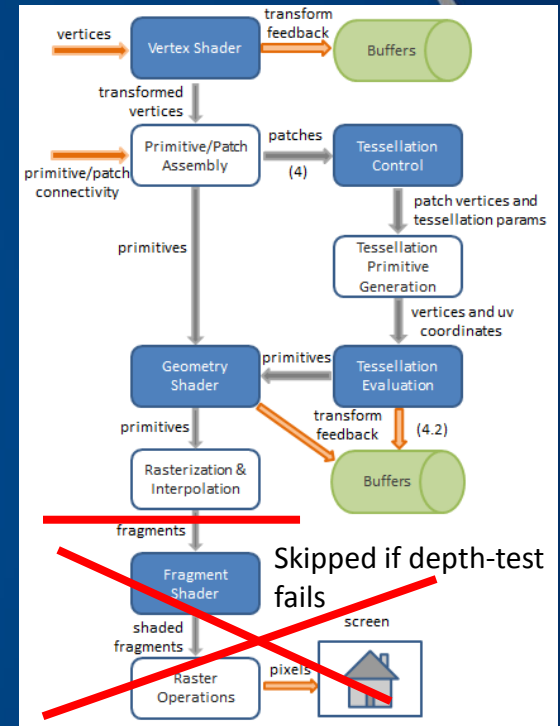
Enter the Z-buffer

- A solution to the render order problem was needed
 - The Z-Buffer was created to solve it
- The Z-Buffer can be thought of as a texture buffer, but instead of pixel colours it holds depth information at each pixel
- This depth information is then used by the Rasteriser
 - Typically any pixel currently being rendered will test its distance against the Z-buffer
 - If the pixel is “behind” the one already rendered then it is discarded




Front-to-Back Ordering

- Using the Z-Buffer we start a frame by clearing it to the furthest distance possible
 - To ensure everything we draw is “in front” of it
- We then render all the meshes and if a mesh renders into a pixel the depth is set
 - An object further away trying to render into the pixel will be caught by the Rasteriser and be ignored
- All geometry for a mesh will still go through the render pipeline but will end before the Fragment Shader
 - Modern shading algorithms are heavily pixel-based, so this can save considerable pixel processing time!
- Highly populated scenes can benefit from rendering the nearest objects first to reduce re-rendering pixels; a problem called Overdraw



Accessing Depth

- With GLSL we can actually access a pixel's depth information from within the Fragment Shader
- There is a built-in global variable, **gl_FragCoord**
 - A **vec4**
 - X and Y refer to the window coordinates of the current pixel fragment
 - Z refers to the depth
 - W actually stores $1/W$ which was used during the projection from 3D space to 2D window coordinates



Summary

- Render order matters!
- A Z-buffer takes care of culling any pixels being drawn behind already drawn pixels
 - Must remember to “wipe” the z-buffer each frame
 - Not clearing the z-buffer can cause some interesting side effects

http://en.wikipedia.org/wiki/Painter's_algorithm

<http://en.wikipedia.org/wiki/Z-buffering>